Bayes R2 for LM

Silvano

April 18, 2018

1   26.03.18

The following R^2 may be useful:

+ $R^2_{Gelman} = \frac{V(\sum_{n=1}^{N} \hat{y}_n^s)}{V(\sum_{n=1}^{N} \hat{y}_n^s) + V(\sum_{n=1}^{N} e_n^s)}$, where $\theta^s, s = 1, ..., S$ are draws from the posterior paramter distribution. For each θ^s we can compute the vector of predicted values $\hat{y}_n^s = E(y \mid X_n, \theta^s)$ and the vector of errors $e_n^s = y_n - \hat{y}_n^s$.

    This R^2 can be interpreted as a data-based estimate of the proportion of variance explained for new data (Gelman). They note that their proposed R^2 overestimates the proportion of explained variance for new data. They argue, a new issue then arises,when fitting a set of a models to a single dataset. Now that the denominator of R^2 is no longer fixed, we can no longer interpret an increase in R^2 as a improved fit to a fixed target.

    Instead of the above error definition we may aswell use samples from the posterior σ_s in the denominator for the linear model. I think this way we would have a better possibility to include prior information in the sigma parameter. In the uninformativ prior case this should not make a difference.

+ In the same paper they note: An alternative is to summarize residual error on the log-probability scale, in which case we recommend using approximate leave-one-out cross-validation (Vehtari, 2017), which can be done using the loo function within rstanarm. If desired, changes in expected log-probability scores can be converted back to an R2 scale as discussed by Nagelkerke (1991).

    instead of leave-one out crossvalidation we may just use the whole data:

    $R^2_{LRT} = 1 - \left(\frac{L_{intercept}}{L_{full}}\right)^{\frac{2}{n}}$, where $L =$ Likelihood or in loglikelihood terms $1 - exp(-LRT/n)$

    For the L_{full} term we can use the vector of predicted values obtained

from the posterior sample. For the $L_{intercept}$ term we can use the maximum likelihood estimate, the mean, of the data at hand. This R^2 is then lower as if we would use the maximum likelihood for the full model. It then includes the uncertainty of the parameter estimates (the higher the sample size, the more certain we can be).

Even without leave-one out crossvalidation this measure is then substantially lower then the R^2_{Gelman} and can be negative. I think that it can be negativ is not a big problem in this case. It measures gemoetric mean improvement per observation for new data. It predicts the mean for new data and can therefor be lower then if we would observe it. Using leave one out crossvalidation with maximum likelihood the R^2 can also be negativ if the correct R^2 formula is used (Beware of R2: simple, unambiguous assessment of the prediction accuracy of QSAR and QSPR models, 2015). A Problem with the LMG formula may arise. However, the non decreasing property when adding predictors is lost in all the bayes formulas described, so this problem is there anyway. A benefit of using the comparison to the fixed mean would be that we are comparing all models to a fixed target.

+ Another option would be to use the classical definition:

$$R^2_{KL} = 1 - \frac{SSE(X)}{SST} \tag{1}$$

$$SSE(X_s) = \sum_{n=1}^{n} (y_i - \hat{y}(x))^2 \tag{2}$$

$$SST_s = \sum_{n=1}^{n} (y_i - \hat{y})^2 \tag{3}$$

$SSE(X_s)$ uses the vector of predicted values from the posterior sample, SST_s uses its sample from the posterior intercept.

Note this R^2 definition is also the Kullback-Leibler divergence R^2 definition for the normal distribution.
This R^2 can be negative and the target is again not fixed.

Some R code giving an example of each case

```
library(INLA); library(brinla); library(xtable); library(mvtnorm)


## Loading required package:  sp

## Loading required package:  Matrix

## This is INLA_17.06.20 built 2017-06-20 03:44:36 UTC.
## See www.r-inla.org/contact-us for how to get help.
```

```
#example data from https://github.com/julianfaraway/brinla

data(usair, package = "brinla")

center_scale <- function(x) {
  scale(x, scale = FALSE)
}
# demean predictors first, need for our R^2_{KL}

usair2<-center_scale(usair[2:7])
SO2<-usair[,1]
usair<-data.frame(cbind(SO2,usair2))

usair.formula1 <-  SO2 ~ negtemp + manuf + wind + precip  + days
usair.lm1 <- lm(usair.formula1, data = usair)
round(summary(usair.lm1)$r.squared,4)
```

[1] 0.604

```
round(summary(usair.lm1)$adj.r.squared,4)
```

[1] 0.5475

```
#R2 from Gelman implemented in rstanarm package using MCMC

library(rstanarm)
```

Loading required package: Rcpp

rstanarm (Version 2.17.2, packaged:)

- Do not expect the default priors to remain the same in future rstanarm versions.

Thus, R scripts should specify priors explicitly, even if they are just the defaults.

- For execution on a local, multicore CPU with excess RAM we recommend calling

options(mc.cores = parallel::detectCores())

- Plotting theme set to bayesplot::theme_default().

```
M1 <- stan_glm(SO2 ~ negtemp + manuf + wind + precip, data=usair )
```

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).

Gradient evaluation took 0.000155 seconds 1000 transitions using 10 leapfrog steps per transition would take 1.55 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration: 1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.092674 seconds (Warm-up) 0.06273 seconds (Sampling) 0.155404 seconds (Total)

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 2).

Gradient evaluation took 1.3e-05 seconds 1000 transitions using 10 leapfrog steps per transition would take 0.13 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration: 1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.095057 seconds (Warm-up) 0.083562 seconds (Sampling) 0.178619 seconds (Total)

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 3).

Gradient evaluation took 1.4e-05 seconds 1000 transitions using 10 leapfrog steps per transition would take 0.14 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration: 1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.092691 seconds (Warm-up) 0.065632 seconds (Sampling) 0.158323 seconds (Total)

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 4).

Gradient evaluation took 1.5e-05 seconds 1000 transitions using 10 leapfrog steps per transition would take 0.15 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration:

1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.100753 seconds (Warm-up) 0.069139 seconds (Sampling) 0.169892 seconds (Total)

```
print(median(bayes_R2(M1)))
```

[1] 0.594538

```
#For the other two R^2 INLA is used.


usair.inla1 <- inla(usair.formula1, data = usair, control.compute = list(dic =
inla.obj<-usair.inla1
samples<-inla.posterior.sample(n=1000, inla.obj)
beta <- matrix(0, nrow = 1000, ncol = 6)
for(i in 1:1000){
  beta[i,] <- samples[[i]]$latent[42:47]
}

yhat <- inla.obj$model.matrix %*% t(beta)

#R2_LRT

M<-matrix(0, 1000,2)
for ( i in 1:1000){

  v<-1/samples[[i]]$hyperpar
  M[i,2]=v
  logLik1 <- dmvnorm(usair$SO2, mean=yhat[,i], sigma=diag(41)*v, log=T)

  #logLik0 <- dmvnorm(usair£SO2, mean=rep(beta[i,1],41), sigma=diag(41)*v, log=
  fit0<-lm(SO2~1, data=usair)
  logLik0 <- logLik(fit0)
  LRT=-2*(logLik0-logLik1)
  M[i,1]=1-exp(-LRT/41)
}
median(M[,1]) # R2_LRT substantially lower
```

[1] 0.5334194

```
# third option classical definition R_KL

M=matrix(0, 1000, 1)
for (i in 1:1000){
M[i,1]=1-var(usair$SO2-yhat[,i])/var((usair$SO2-yhat[1,i]))
}

median(M)
```

[1] 0.5565833

```
#4th option like Gelman but error also sampled from posterior

M=matrix(0, 1000, 1)
for (i in 1:1000){
  M[i,1]=var(yhat[,i])/(var(yhat[,i])+(1/samples[[i]]$hyperpar))
}

median(M) # does not make a big difference
```

[1] 0.5993437

```
          # when prior is uninformativ compared to R^2_Gelman
```

Some Problems with predictors that have no explanatory power

```
set.seed(3)

usair$x1<-scale(rnorm(41,0,1))
usair$x2<-scale(rnorm(41,0,1))
usair$x3<-scale(rnorm(41,0,1))

usair.formula1 <-  SO2 ~ x1 + x2 + x3
usair.lm1 <- lm(usair.formula1, data = usair)
round(summary(usair.lm1)$r.squared,4)
```

[1] 0.0493

```
round(summary(usair.lm1)$adj.r.squared,4)
```

[1] -0.0278

```
M1 <- stan_glm(SO2 ~ x1 + x2 + x3, data=usair )
```

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 1).

Gradient evaluation took 2.4e-05 seconds 1000 transitions using 10 leapfrog steps per transition would take 0.24 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration: 1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.087608 seconds (Warm-up) 0.068254 seconds (Sampling) 0.155862 seconds (Total)

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 2).

Gradient evaluation took 1.5e-05 seconds 1000 transitions using 10 leapfrog steps per transition would take 0.15 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration: 1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.097296 seconds (Warm-up) 0.084512 seconds (Sampling) 0.181808 seconds (Total)

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 3).

Gradient evaluation took 1.3e-05 seconds 1000 transitions using 10 leapfrog steps per transition would take 0.13 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration: 1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.08062 seconds (Warm-up) 0.06521 seconds (Sampling) 0.14583 seconds (Total)

SAMPLING FOR MODEL 'continuous' NOW (CHAIN 4).

Gradient evaluation took 1.2e-05 seconds 1000 transitions using 10 leapfrog steps per transition would take 0.12 seconds. Adjust your expectations accordingly!

Iteration: 1 / 2000 [0Iteration: 200 / 2000 [10Iteration: 400 / 2000 [20Iteration: 600 / 2000 [30Iteration: 800 / 2000 [40Iteration: 1000 / 2000 [50Iteration: 1001 / 2000 [50Iteration: 1200 / 2000 [60Iteration: 1400 / 2000 [70Iteration: 1600 / 2000 [80Iteration: 1800 / 2000 [90Iteration: 2000 / 2000 [100

Elapsed Time: 0.085711 seconds (Warm-up) 0.066193 seconds (Sampling) 0.151904 seconds (Total)

```r
print(median(bayes_R2(M1))) # bigger than maxiumum likelihood R^2
```

[1] 0.09407575

```r
usair.inla1 <- inla(usair.formula1, data = usair, control.compute = list(dic =
inla.obj<-usair.inla1
samples<-inla.posterior.sample(n=1000, inla.obj)
beta <- matrix(0, nrow = 1000, ncol = 4)
for(i in 1:1000){
  beta[i,] <- samples[[i]]$latent[42:45]
}

yhat <- inla.obj$model.matrix %*% t(beta)

#R2_LRT

M<-matrix(0, 1000,2)
for ( i in 1:1000){

  v<-1/samples[[i]]$hyperpar
  M[i,2]=v
  logLik1 <- dmvnorm(usair$SO2, mean=yhat[,i], sigma=diag(41)*v, log=T)

  fit0<-lm(SO2~1, data=usair)
  logLik0 <- logLik(fit0)
  LRT=-2*(logLik0-logLik1)
  M[i,1]=1-exp(-LRT/41)
}
median(M[,1]) # R2_LRT substantially lower
```

[1] -0.05644987

```r
# third option classical definition R_KL

M=matrix(0, 1000, 1)
```

```
for (i in 1:1000){
M[i,1]=1-var(usair$SO2-yhat[,i])/var((usair$SO2-yhat[1,i]))
}

median(M)
```

[1] -0.00889446

```
#4th option like Gelman but error also sampled from posterior

M=matrix(0, 1000, 1)
for (i in 1:1000){
  M[i,1]=var(yhat[,i])/(var(yhat[,i])+(1/samples[[i]]$hyperpar))
}

median(M)
```

[1] 0.09052242

The R^2_{Gelman} is even bigger than the R^2 from the lm object using maximum likelihood. Comparing Baysian R^2 : $R^2_{Gelman} > R^2_{KL} > R^2_{LRT}$, which was expected.